

serverless architectures



functions, events, cloud native services and things like that...

Dave Townsend

Principal Software Engineer
Innovation & Architecture
Matson, Inc

@davetownsend




what is serverless?

O'REILLY®

What is Serverless?

Understanding the Latest Advances in
Cloud and Service-Based Architecture



Mike Roberts
& John Chapin

*“**Cloud** really just means **datacenterless**, there’s a datacenter somewhere but we don’t care.”*

– Adrian Cockcroft

serverless

we don't have to think about the servers

also...

A Serverless solution is one that costs you nothing to run if nobody is using it ... excluding data storage costs.

– Paul Johnston

why serverless?

(brief) history of compute

(brief) history of compute

physical servers

(brief) history of compute

VMs (on-prem → cloud)

physical servers

(brief) history of compute

containers

VMs (on-prem → cloud)

physical servers

(brief) history of compute

serverless

containers

VMs (on-prem → cloud)

physical servers

(brief) history of compute

serverless

containers



VMs (on-prem → cloud)

physical servers

(brief) history of compute

serverless

value line

containers



VMs (on-prem → cloud)

physical servers

instances are not *really* cloud

instances are not *really* cloud

instances do not take advantage of the cloud

instances are not *really* cloud

instances do not take advantage of the cloud

use the services, the services are good



S3 was serverless before serverless was cool.



DynamoDB



SNS



Athena



Lambda

serverless is about
using **managed services**



S3



Kinesis



API Gateway

next evolution in cloud

backend as-a service & functions as-a service

BaaS

FaaS



Auth0



Firebase



netlify



AWS Lambda



Azure Functions




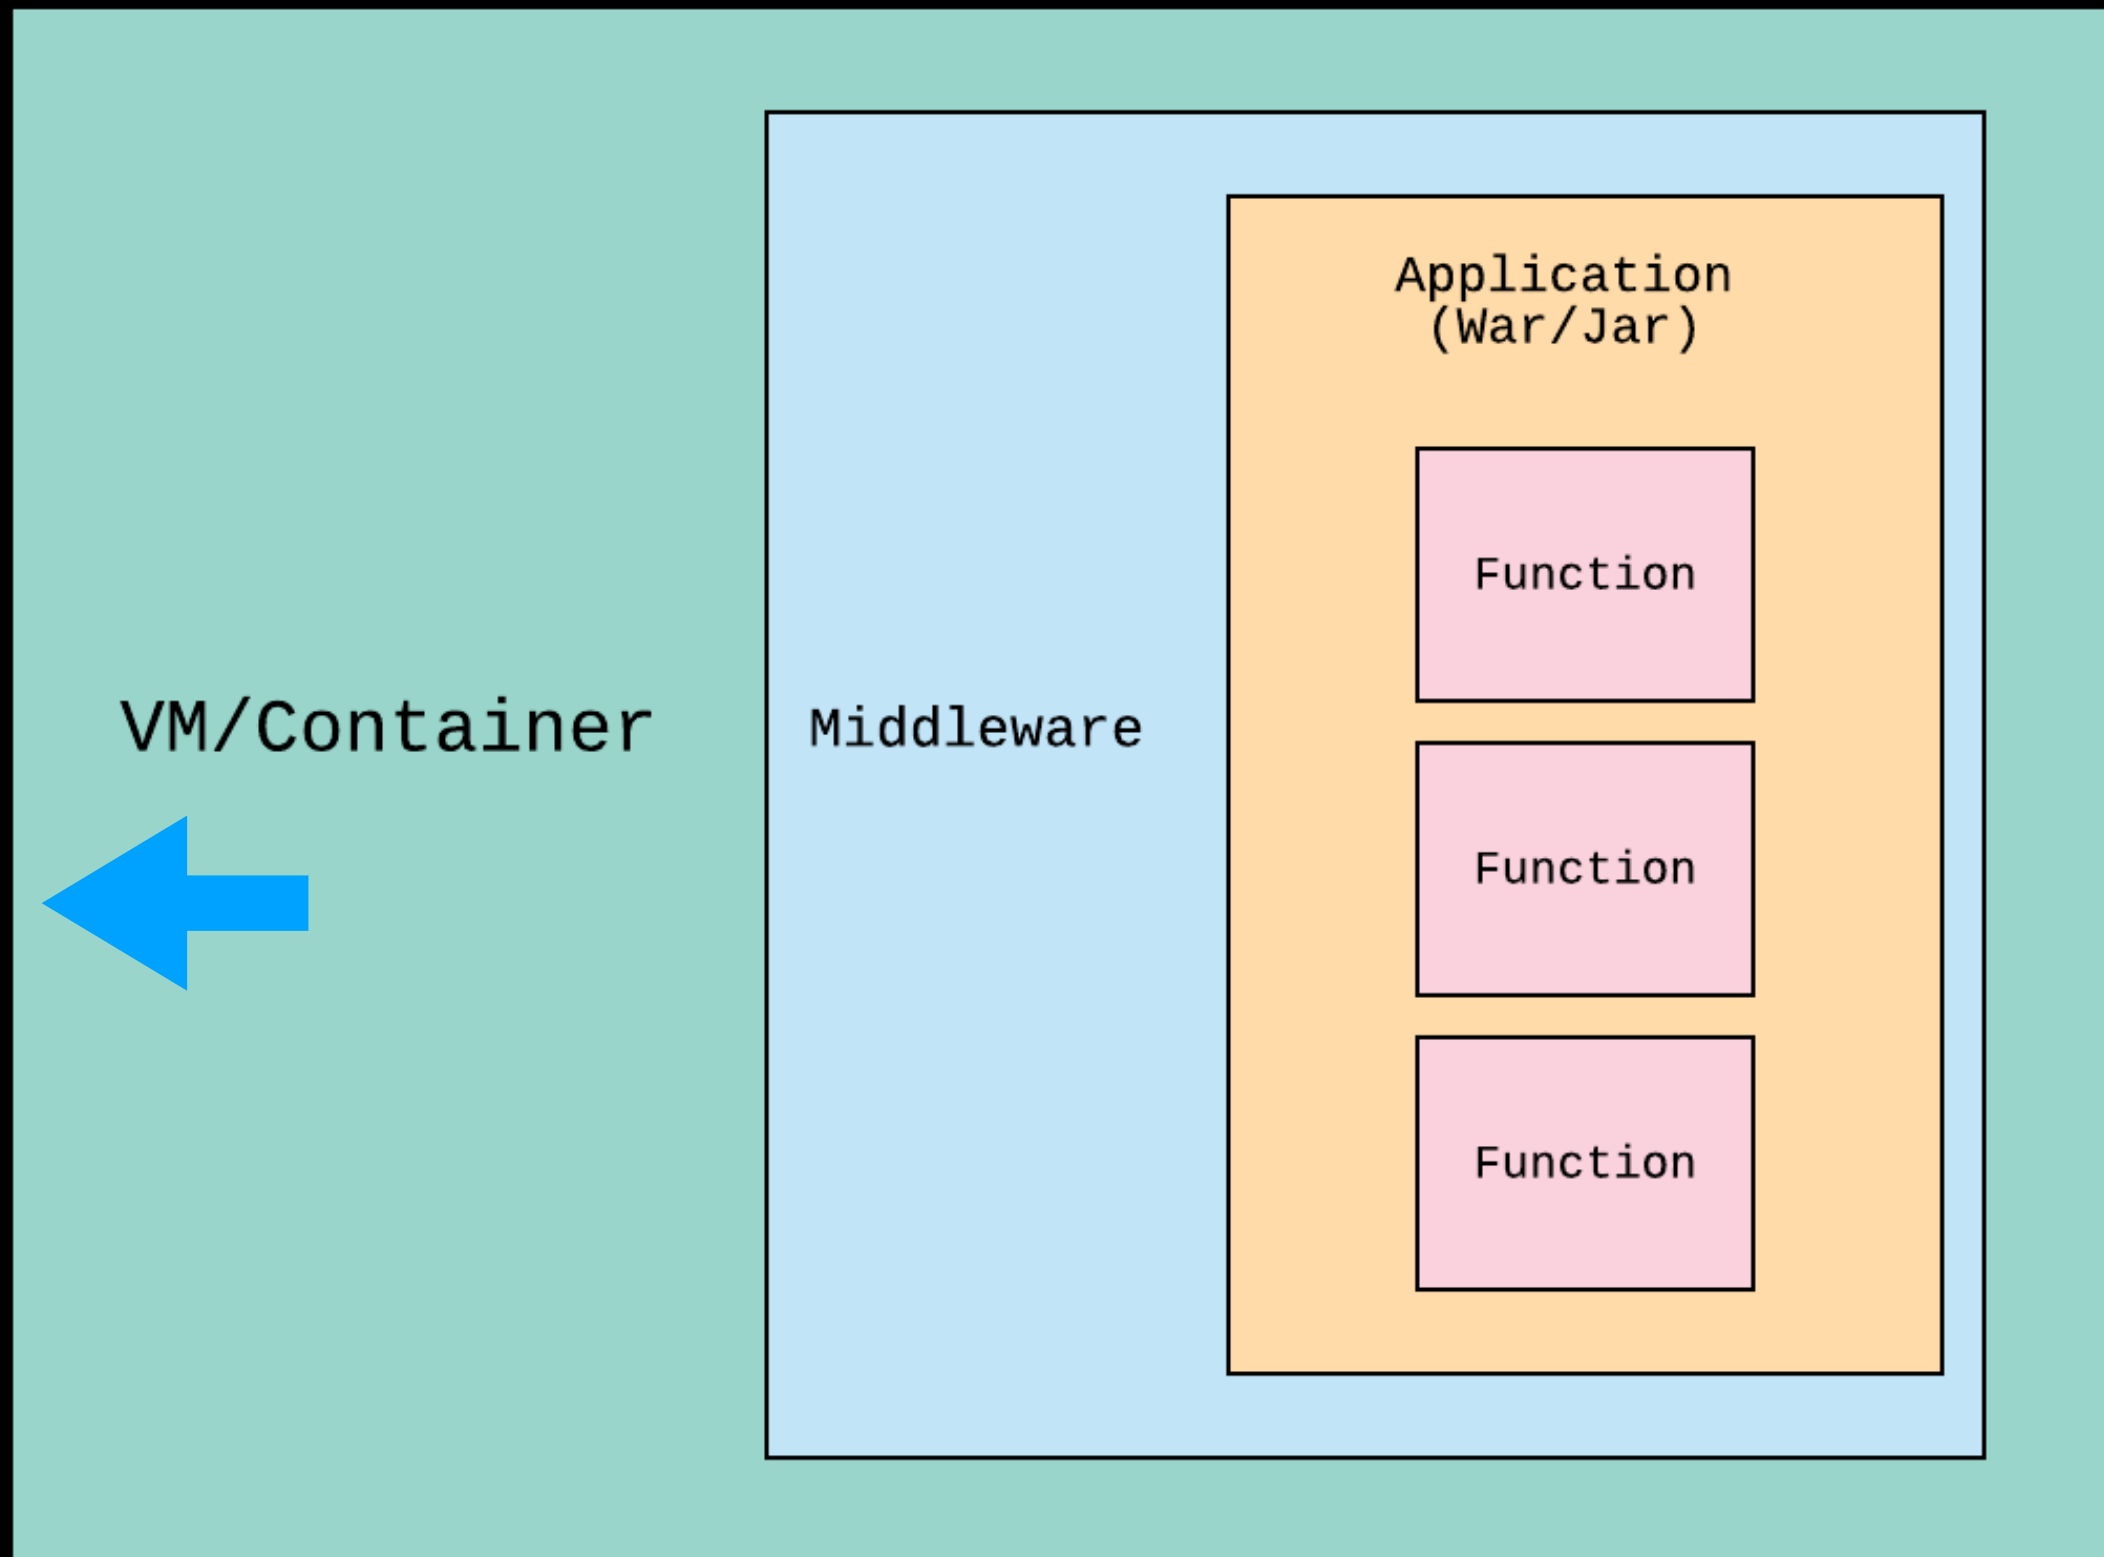
IBM OpenWhisk



Google Cloud Functions


FaaS







VM/Container



Middleware

Application
(War/Jar)

Function

Function

Function

FaaS Provider

Function

Function

Function

FaaS Provider

idle


Function

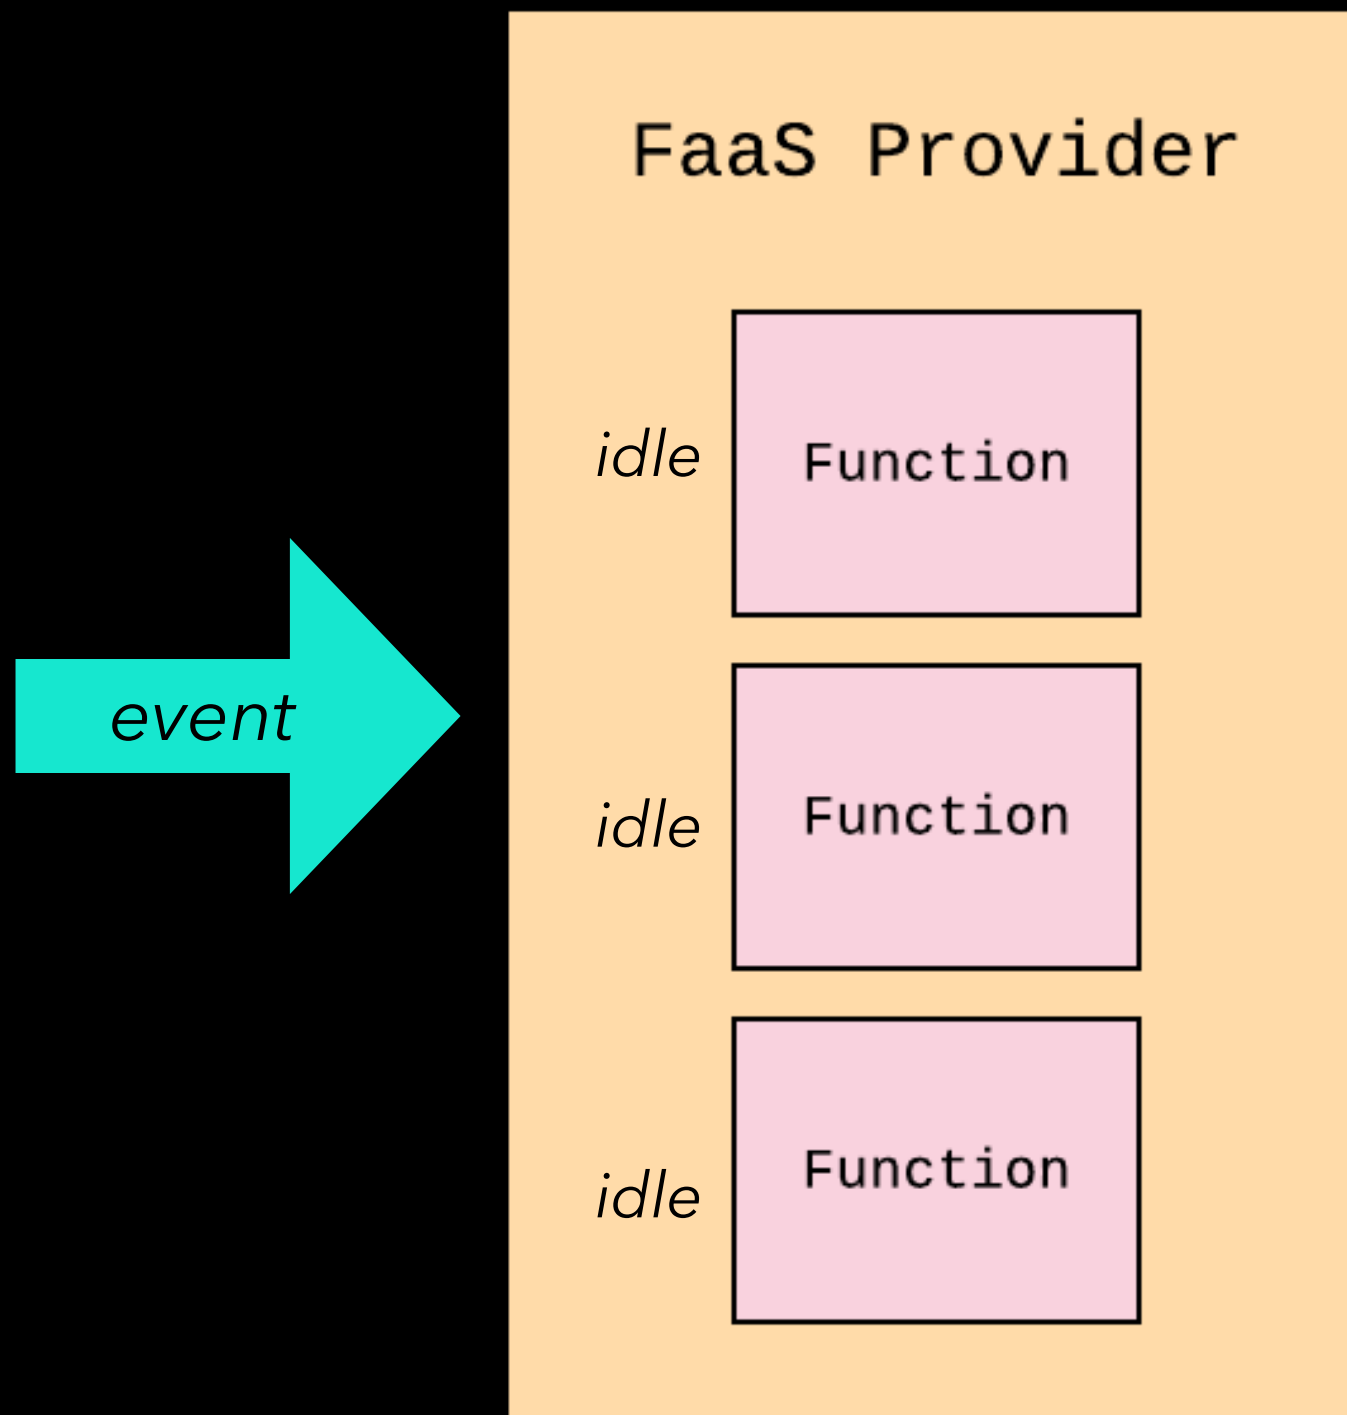
idle

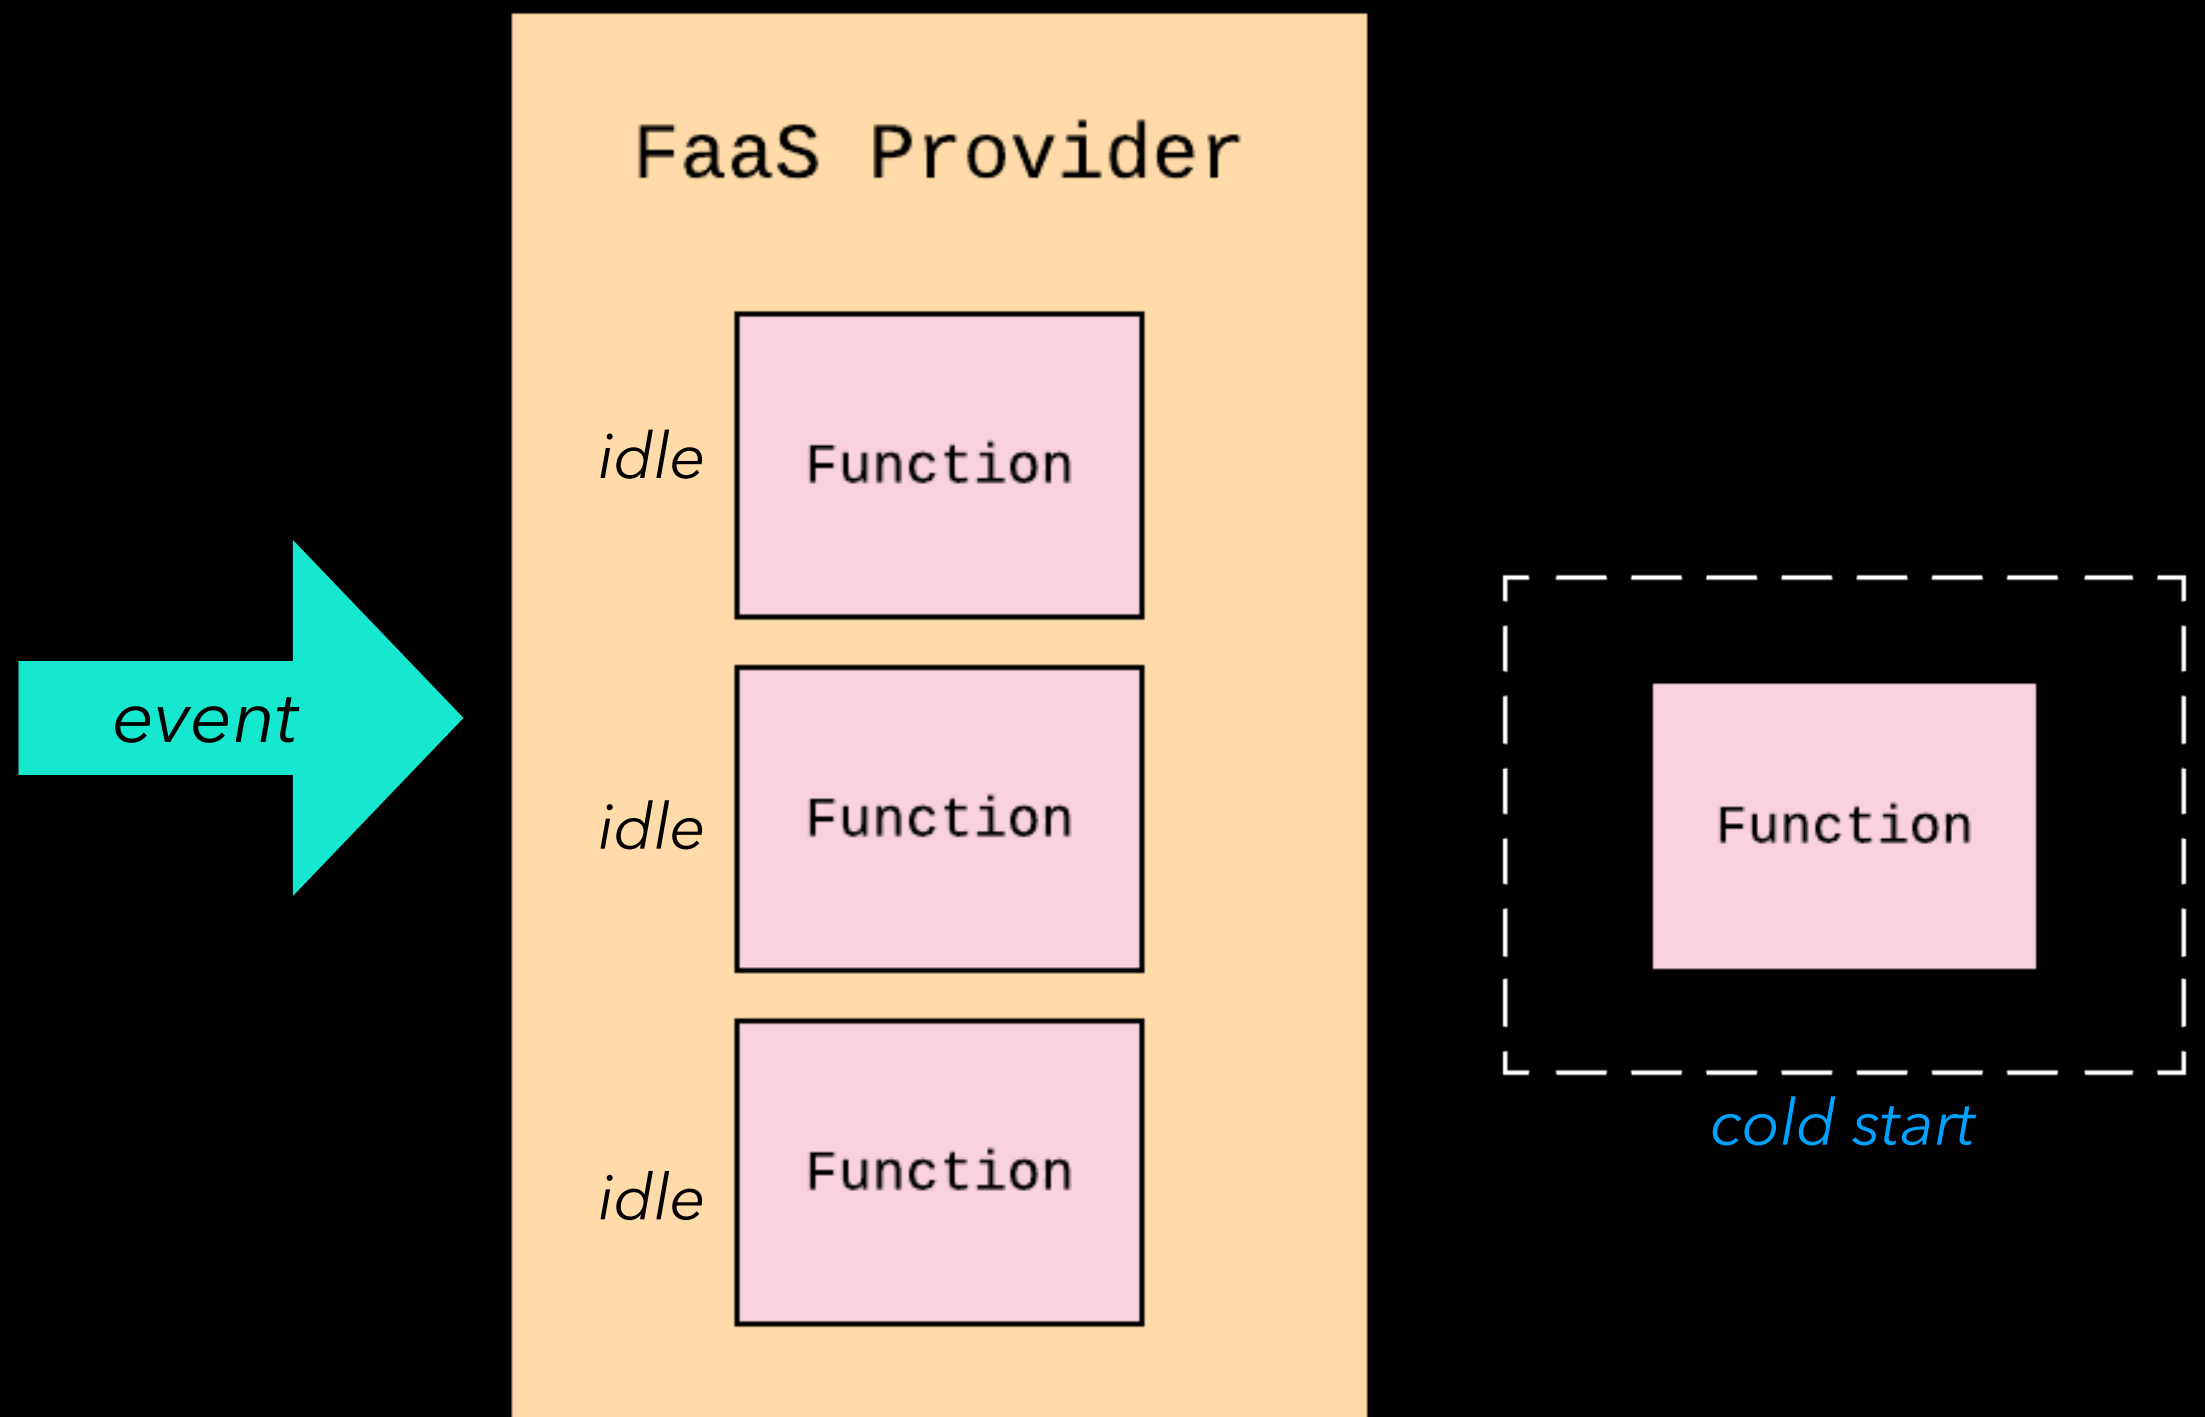
Function

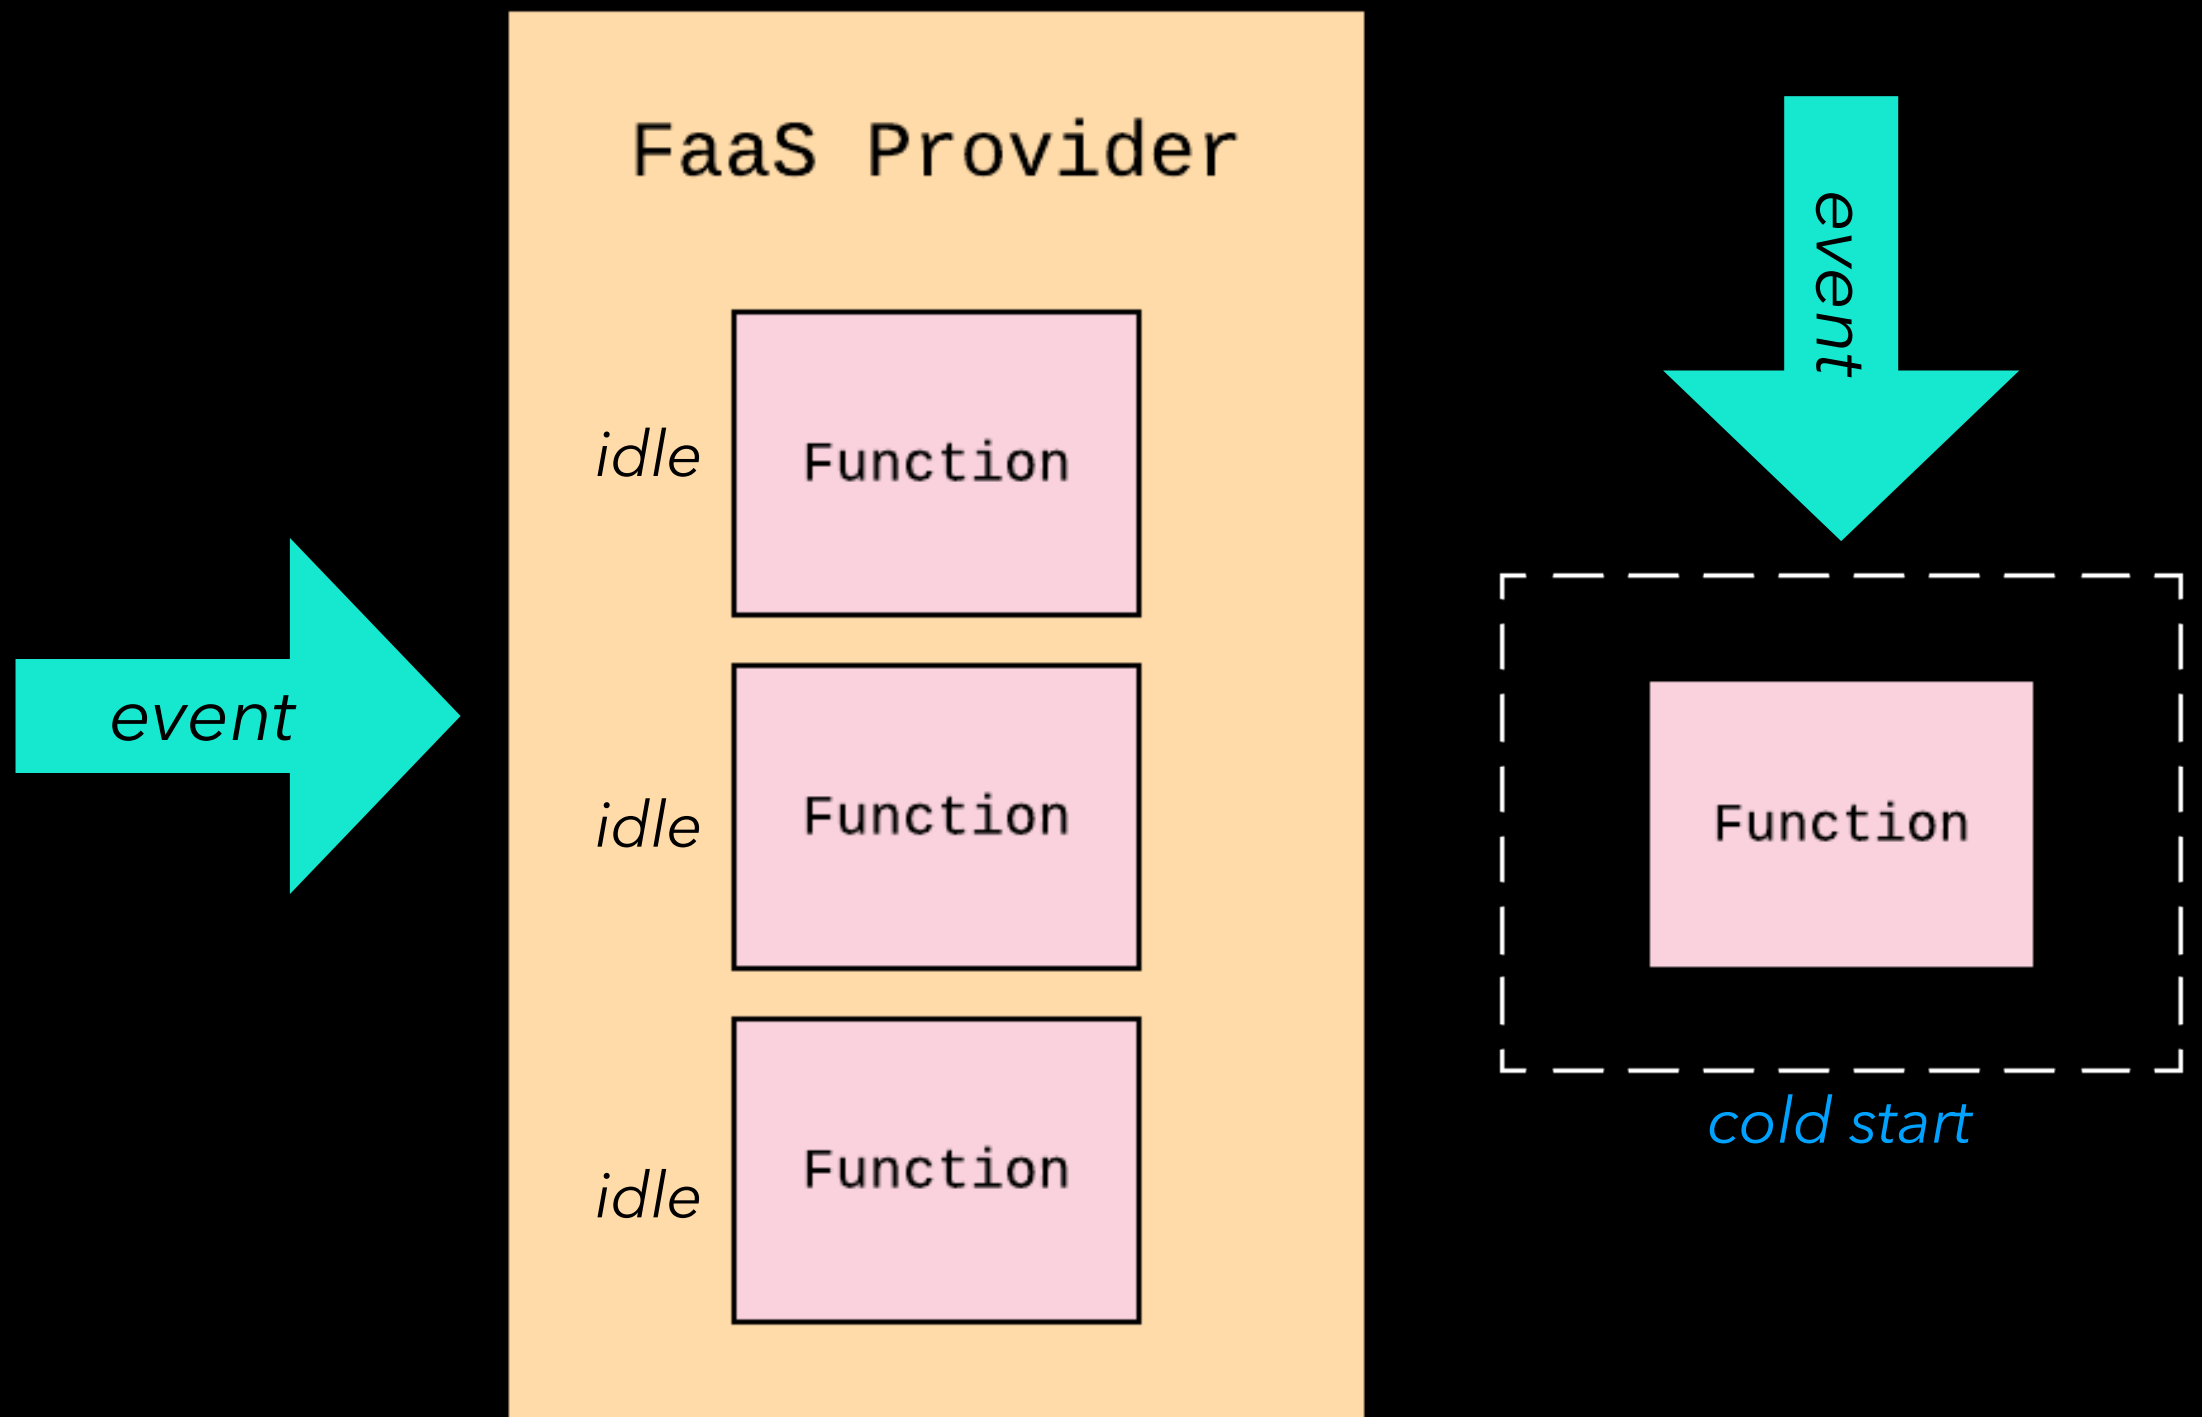
idle

Function









FaaS Provider

idle


Function

idle

Function

idle

Function



lambda language runtimes

Node.js

Java*

Python

C#

lambda language runtimes

Node.js

Java*

Python

C#

framework d'jour not needed here

lambda pricing

first 1 million requests per month are free

\$0.20 per 1 million requests thereafter

\$0.0000002 per request

(billed in 100ms increments)

lambda runtime config

(23 memory choices) 128MB - 1.5GB

timeout (1- 300 sec)

app logging → CloudWatch

- cwtail!

/tmp gotcha!

serverless characteristics

...not limited to functions.

implicit scalability & HA

built in autoscale and provisioning

implicit HA by default


does *not* mean implicit DR

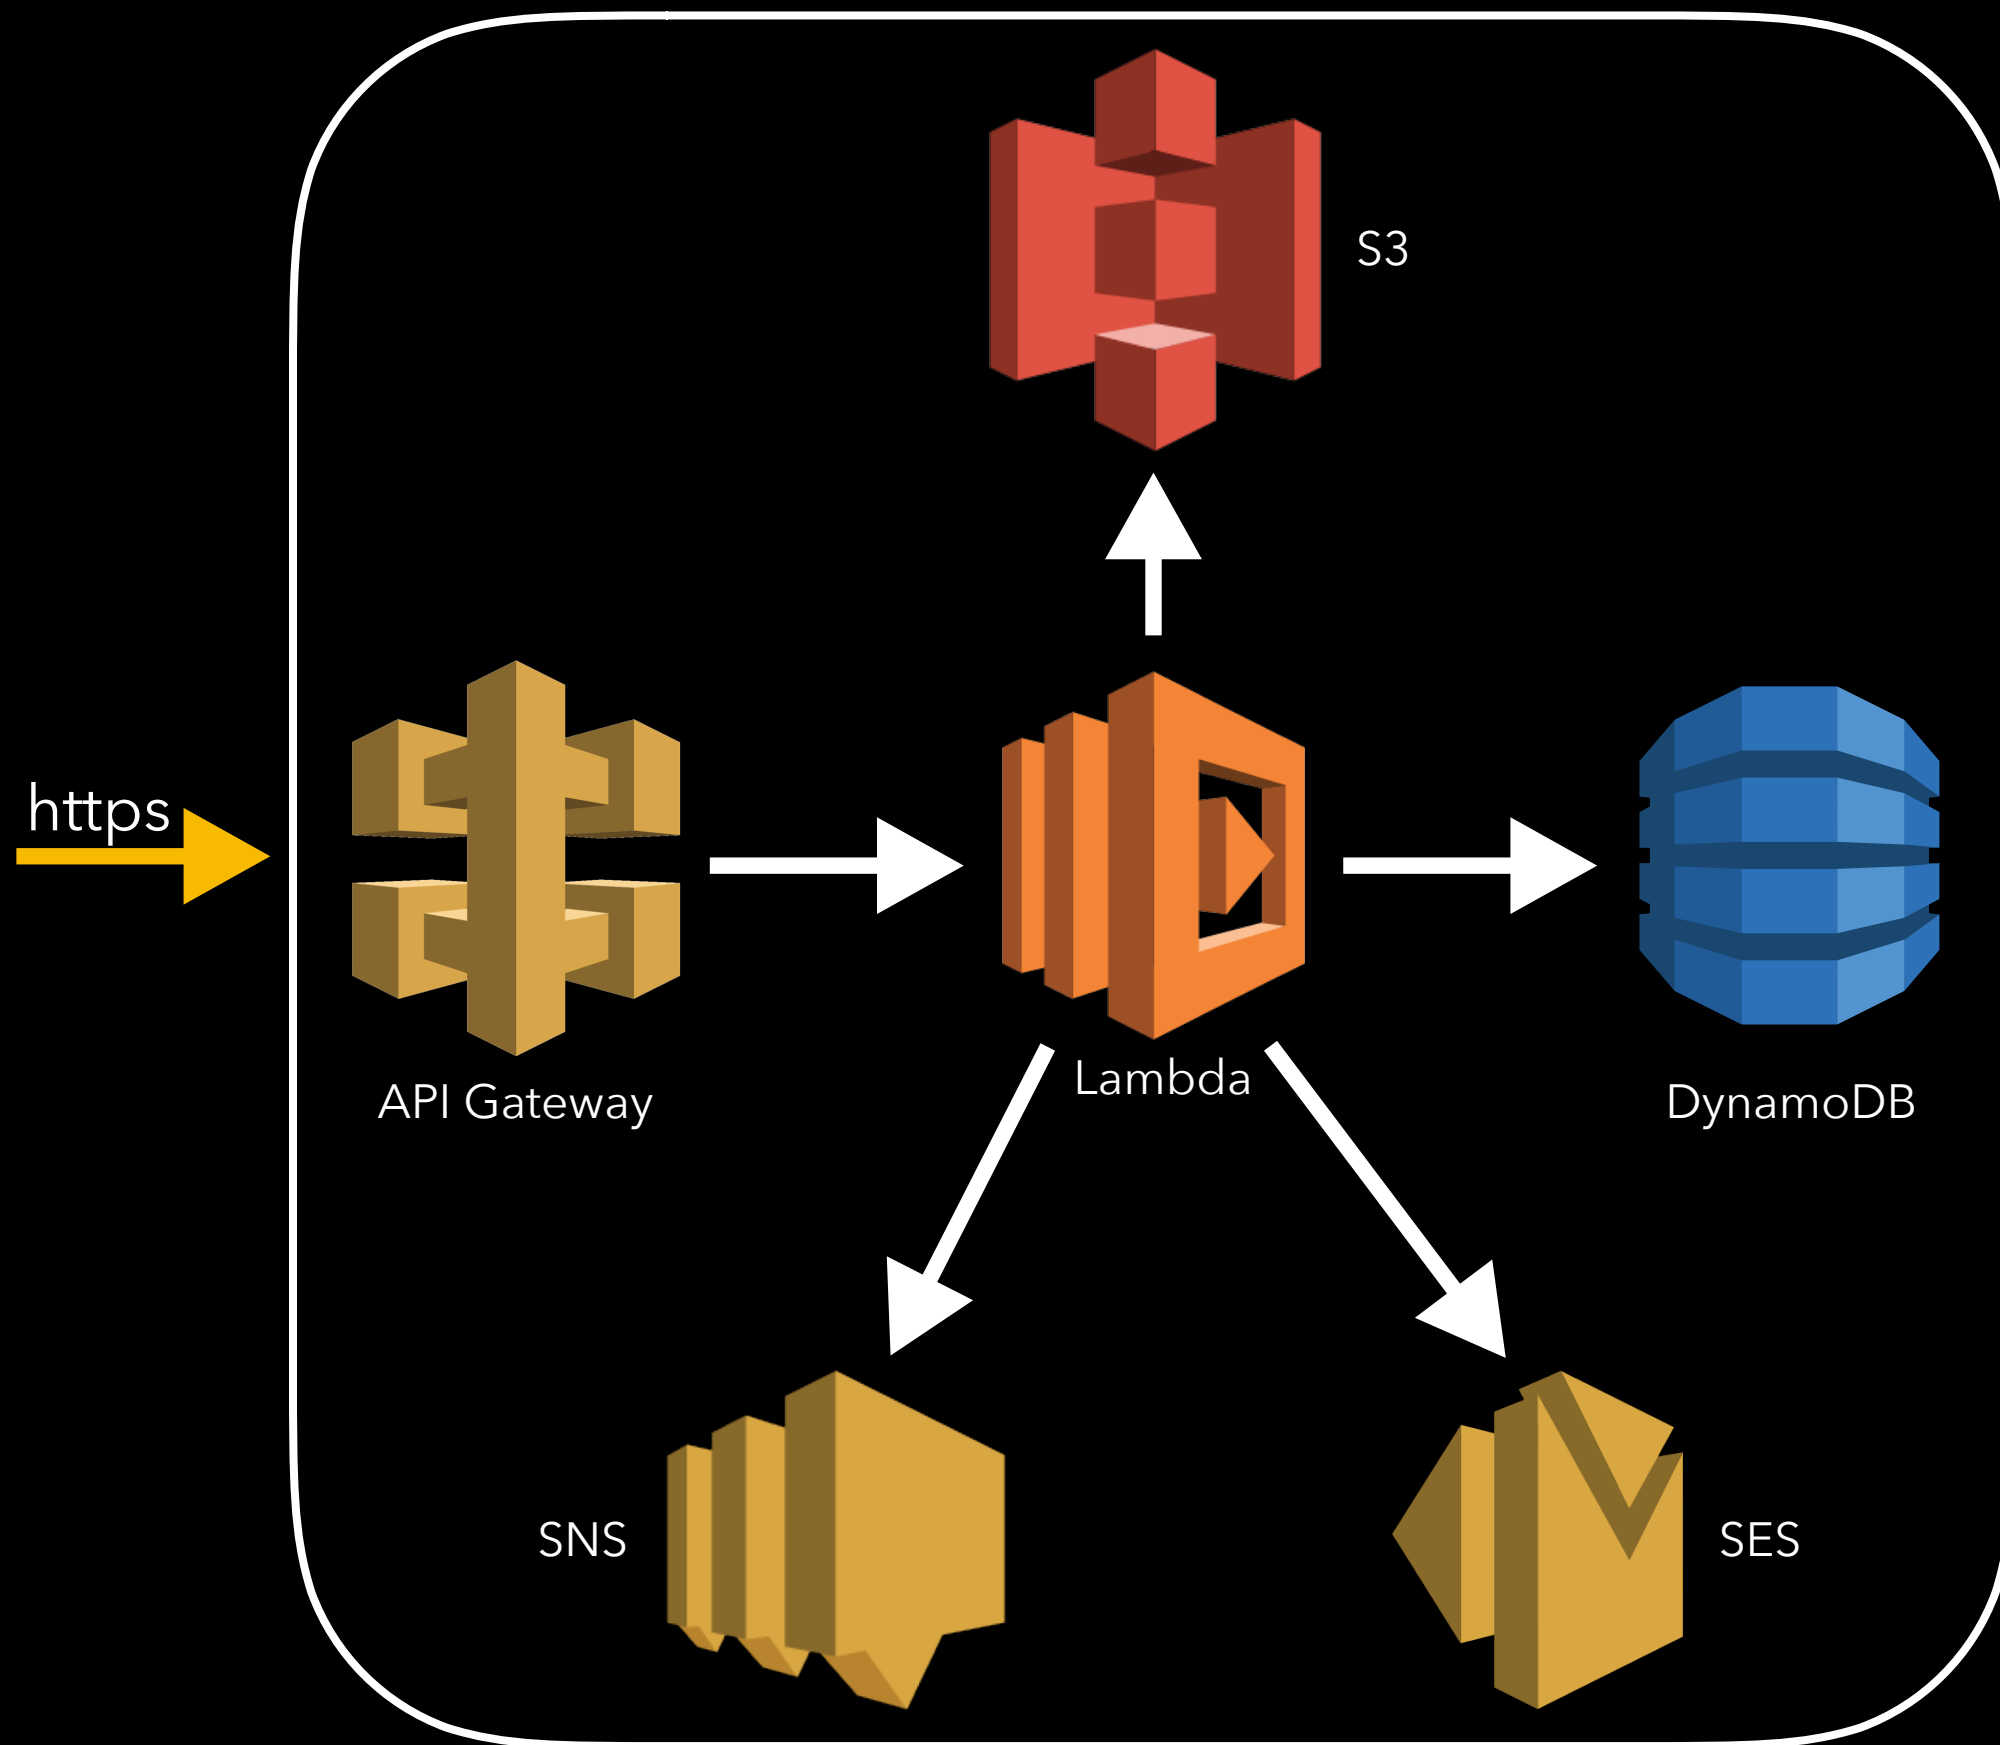
pay-per execution

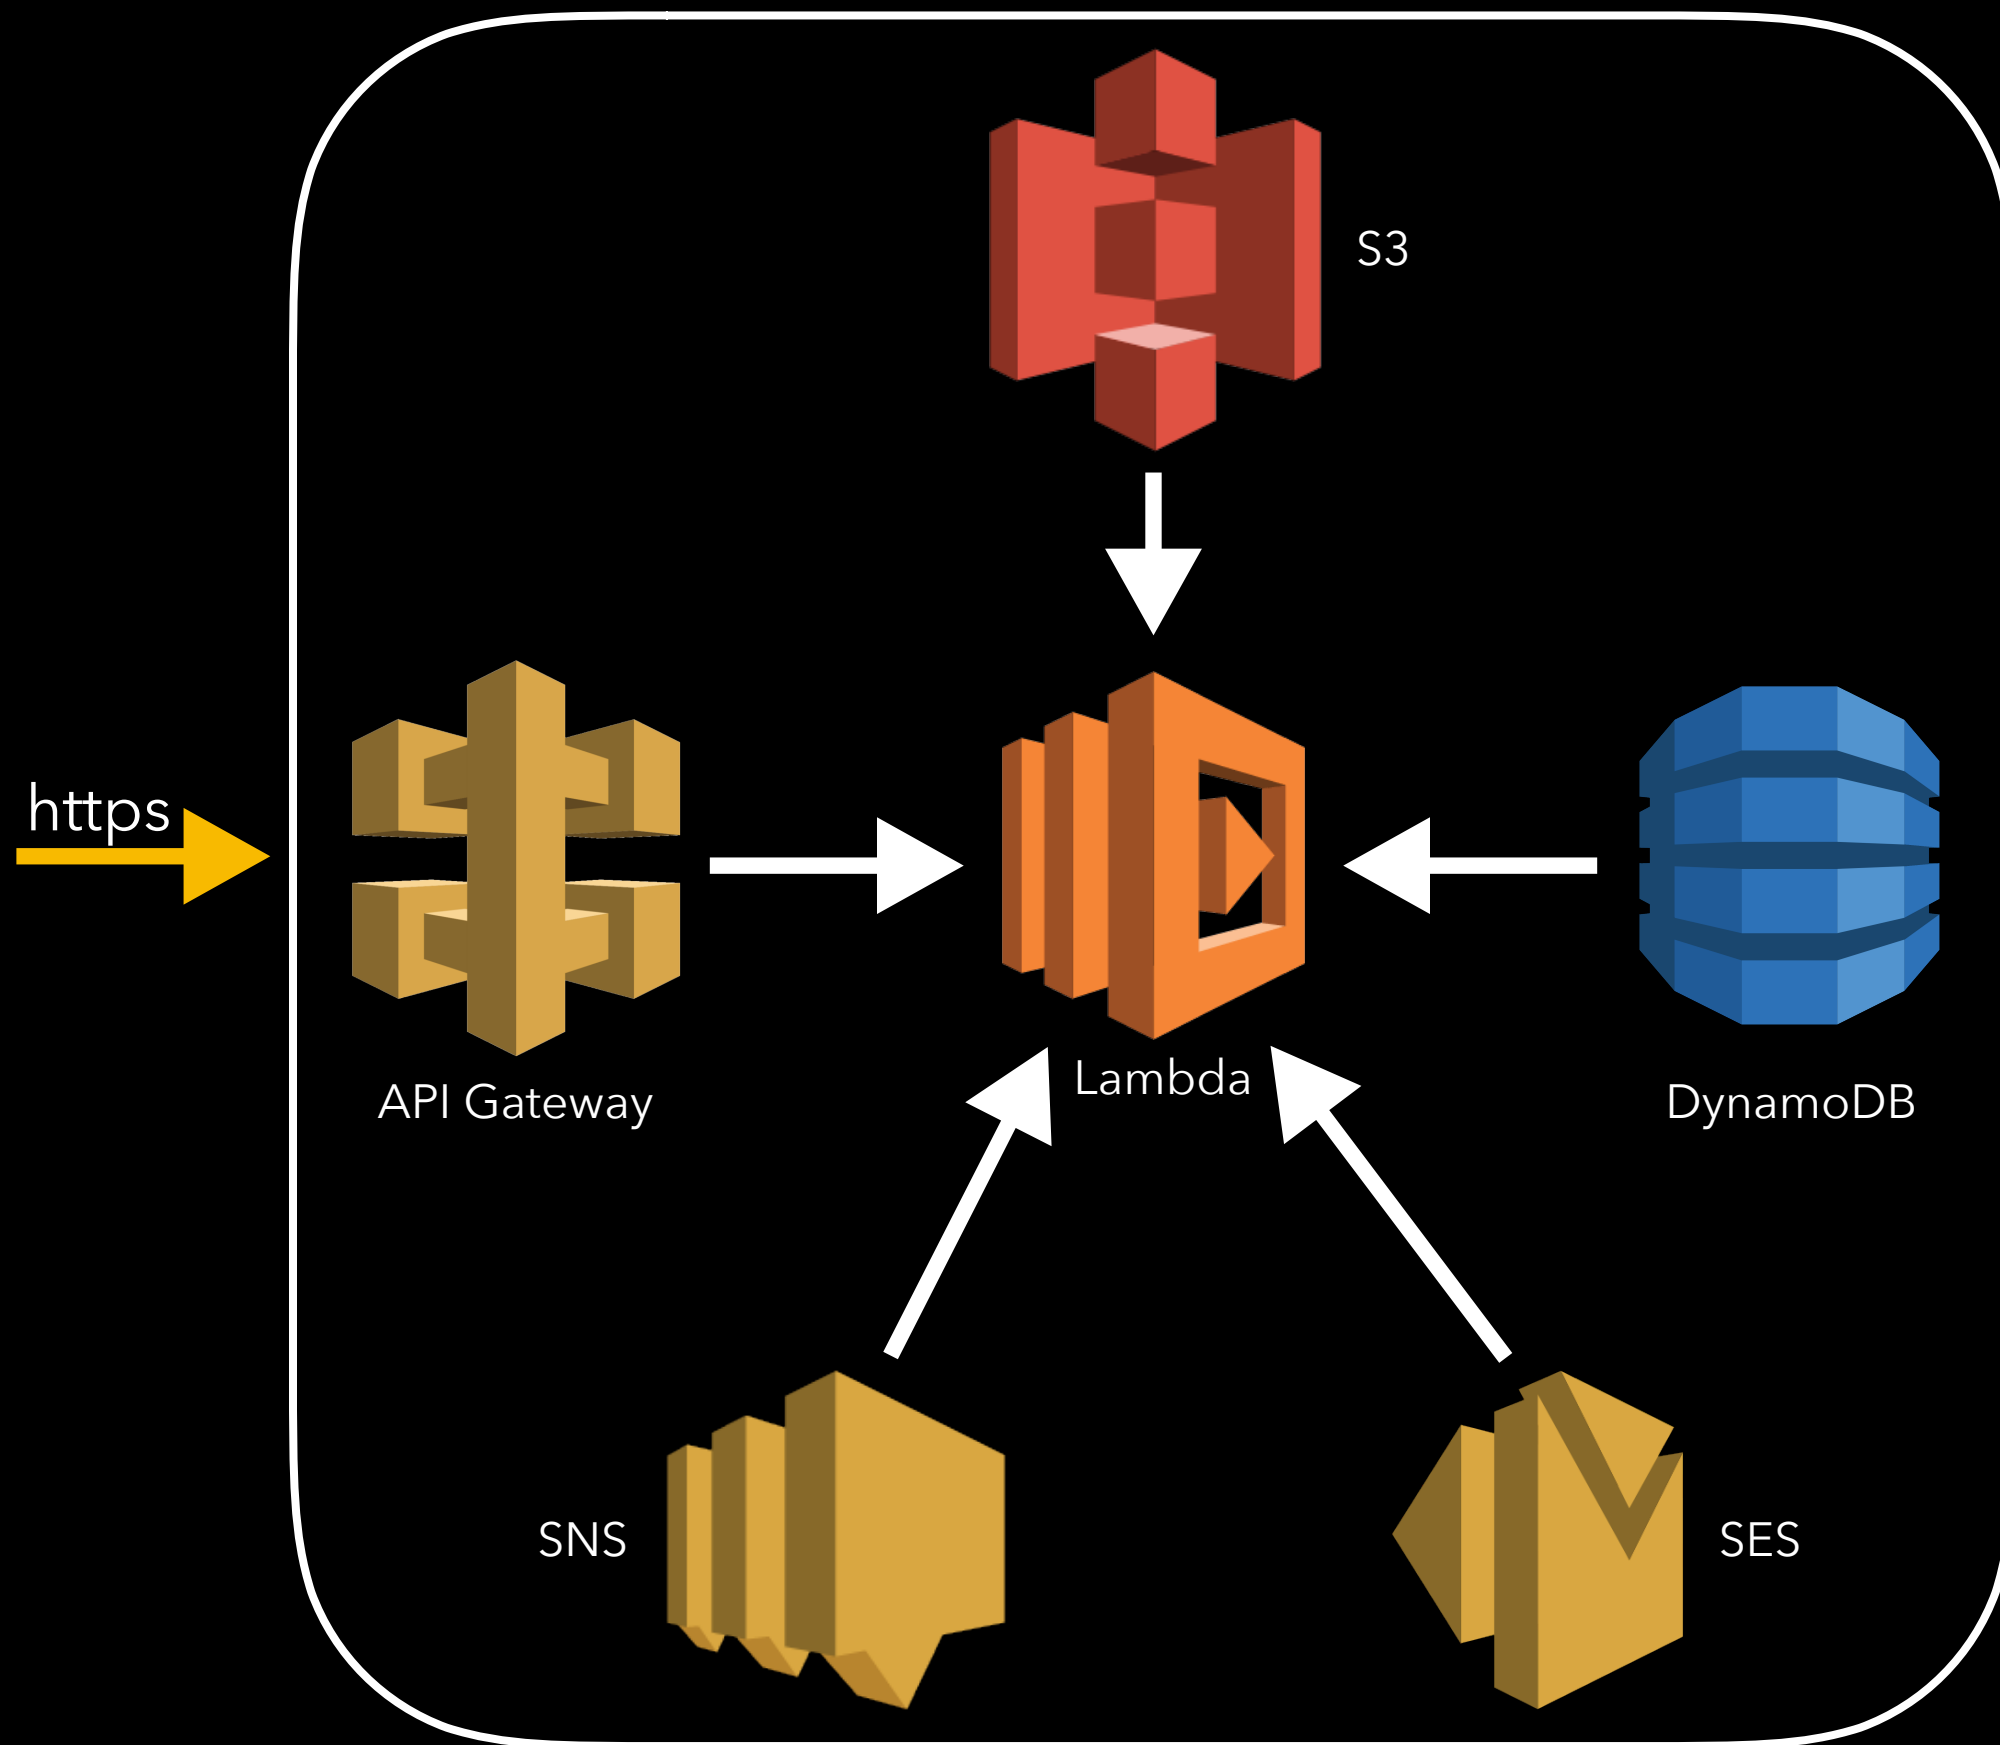
never pay for idle

billed for usage only

no/low-cost DR 😊👍







- Amazon S3
- Amazon DynamoDB
- Amazon Kinesis Streams
- Amazon Simple Notification Service
- Amazon Simple Email Service
- Amazon Cognito
- AWS CloudFormation
- Amazon CloudWatch Logs
- Amazon CloudWatch Events
- AWS CodeCommit
- Scheduled Events (powered by Amazon CloudWatch Events)
- AWS Config
- Amazon Alexa
- Amazon Lex
- Amazon API Gateway

current list (Sept 2017) of supported event sources

“Event-driven architecture will become an essential skill in supporting digital business transformation By 2018”

- Gartner

<http://www.gartner.com/newsroom/id/3758164>

use cases (work loads)

use cases (work loads)

not for everything....

good for *a lot* of things

questionable work loads

long running batch processes

- *lambda max run-time 5min*

low-latency systems

heavy continuous cpu load

- *~ 75% is break even on lambda cost*

popular work loads

web applications

mobile backends

data stream processing

chat bots

IoT

alexa skills


OPS automation



security auditing





video transcoding


microservices


microservices




















serverless delivers the original
promise of the cloud

serverless delivers the original
promise of the cloud

cloud native by default...
unless you really mess things up

serverless awareness

stateless (yes/no)

latency (cold starts)

local testing

debugging

tooling

db's (scale)

security

security

granular IAM policies (key!)

no long standing servers

code dependencies

data

vendor lock-in

vendor lock-in


services and data (comes with territory)



no lock-in is a myth













keep function logic in libs (lambda as wrappers)

cloud agnostic not practical

pick a solid provider and exploit it!









ServerLess Framework



ServerLess Framework

provisioning
orchestration
&
deployment



ServerLess Framework

CloudFormation abstraction

+

aws cli

```
$ npm -i serverless -g
```

```
$ npm -i serverless -g
```

```
$ sls create --template aws-nodejs --path serverless-demo
```



```
$ npm -i serverless -g
```

```
$ sls create --template aws-nodejs --path serverless-demo
```

```
$ sls create --template aws-java-gradle --path serverless-demo
```

serverless.yml x

```
3  provider:
4    name: aws
5    region: us-west-2
6    runtime: nodejs6.10
7
8  functions:
9    trackShipmentByContainer:
10     handler: containerTrack.trackShipmentByContainer
11     memorySize: 256
12     timeout: 10
13     events:
14     - http:
15       path: /shipments/cn/{containerNumber}
16       method: get
17       cors: true
18
19     forwardSNSToSlack:
20     handler: slackForwarder.forwardSNSToSlack
21     memorySize: 256
22     timeout: 10
23     events:
24     - sns:
25       topicName: track-my-container-alerts
26       displayName: Track My Container Alerts
27
28  resources:
29    Resources:
30      DynamoDbTable:
31        Type: AWS::DynamoDB::Table
32        Properties:
33          TableName: tmc-config
34          AttributeDefinitions:
35          - AttributeName: configKey
36            AttributeType: S
37          KeySchema:
38          - AttributeName: configKey
39            KeyType: HASH
40          ProvisionedThroughput:
41            ReadCapacityUnits: 2
42            WriteCapacityUnits: 1
```

```
3 provider:
4   name: aws
5   region: us-west-2
6   runtime: nodejs6.10
7
8 functions:
9   trackShipmentByContainer:
10    handler: containerTrack.trackShipmentByContainer
11    memorySize: 256
12    timeout: 10
13    events:
14      - http:
15        path: /shipments/cn/{containerNumber}
16        method: get
17        cors: true
18
19   forwardSNSToSlack:
20    handler: slackForwarder.forwardSNSToSlack
21    memorySize: 256
22    timeout: 10
23    events:
24      - sns:
25        topicName: track-my-container-alerts
26        displayName: Track My Container Alerts
27
28 resources:
29   Resources:
30     DynamoDbTable:
31       Type: AWS::DynamoDB::Table
32       Properties:
33         TableName: tmc-config
34         AttributeDefinitions:
```

```

3 provider:
4   name: aws
5   region: us-west-2
6   runtime: nodejs6.10
7
8 functions:
9   trackShipmentByContainer:
10    handler: containerTrack.trackShipmentByContainer
11    memorySize: 256
12    timeout: 10
13    events:
14      - http:
15        path: /shipments/cn/{containerNumber}
16        method: get
17        cors: true
18
19    forwardSNSToSlack:
20      handler: slackForwarder.forwardSNSToSlack
21      memorySize: 256
22      timeout: 10
23      events:
24        - sns:
25          topicName: track-my-container-alerts
26          displayName: Track My Container Alerts
27
28 resources:
29   Resources:
30     DynamoDbTable:
31       Type: AWS::DynamoDB::Table
32       Properties:
33         TableName: tmc-config
34         AttributeDefinitions:
35           - {AttributeName: containerNumber, AttributeType: S}

```

```

3 provider:
4   name: aws
5   region: us-west-2
6   runtime: nodejs6.10
7
8 functions:
9   trackShipmentByContainer:
10    handler: containerTrack.trackShipmentByContainer
11    memorySize: 256
12    timeout: 10
13    events:
14      - http:
15        path: /shipments/cn/{containerNumber}
16        method: get
17        cors: true
18
19   forwardSNSToSlack:
20    handler: slackForwarder.forwardSNSToSlack
21    memorySize: 256
22    timeout: 10
23    events:
24      - sns:
25        topicName: track-my-container-alerts
26        displayName: Track My Container Alerts
27
28 resources:
29   Resources:
30     DynamoDbTable:
31       Type: AWS::DynamoDB::Table
32       Properties:
33         TableName: tmc-config
34         AttributeDefinitions:
35           - {AttributeName: containerNumber, AttributeType: S}

```

```
3 provider:
4   name: aws
5   region: us-west-2
6   runtime: nodejs6.10
7
8 functions:
9   trackShipmentByContainer:
10    handler: containerTrack.trackShipmentByContainer
11    memorySize: 256
12    timeout: 10
13    events:
14      - http:
15        path: /shipments/cn/{containerNumber}
16        method: get
17        cors: true
18
19   forwardSNSToSlack:
20    handler: slackForwarder.forwardSNSToSlack
21    memorySize: 256
22    timeout: 10
23    events:
24      - sns:
25        topicName: track-my-container-alerts
26        displayName: Track My Container Alerts
27
28 resources:
29   Resources:
30     DynamoDbTable:
31       Type: AWS::DynamoDB::Table
32       Properties:
33         TableName: tmc-config
34         AttributeDefinitions:
```

```

12     timeout: 10
13     events:
14         - http:
15             path: /shipments/cn/{containerNumber}
16             method: get
17             cors: true
18
19     forwardSNSToSlack:
20         handler: slackForwarder.forwardSNSToSlack
21         memorySize: 256
22         timeout: 10
23         events:
24             - sns:
25                 topicName: track-my-container-alerts
26                 displayName: Track My Container Alerts
27
28 resources:
29     Resources:
30         DynamoDbTable:
31             Type: AWS::DynamoDB::Table
32             Properties:
33                 TableName: tmc-config
34                 AttributeDefinitions:
35                     - AttributeName: configKey
36                       AttributeType: S
37                 KeySchema:
38                     - AttributeName: configKey
39                       KeyType: HASH
40                 ProvisionedThroughput:
41                     ReadCapacityUnits: 2
42                     WriteCapacityUnits: 1

```

deployment

deployment

```
$ sls deploy
```

serverless deployments are
infrastructure *and* code

serverless deployments are
infrastructure *and* code

infrastructure as code 

areas to explore

observability

serverless ops

hybrid serverless applications

resources

Serverless Architectures

<https://martinfowler.com/articles/serverless.html>

The Future of Serverless Compute

<https://www.infoq.com/articles/future-serverless>

Why the Fuss About Serverless

<https://hackernoon.com/why-the-fuss-about-serverless-4370b1596da0>

Serverless Security: What's Left To Protect

https://youtu.be/CiyUD_rI8D8

Serverless Computing and Applications

<https://aws.amazon.com/serverless>

Serverless guide

<https://serverless.github.io/guide>

Serverless Framework

<https://serverless.com/framework>

